



REVA
UNIVERSITY
BENGALURU INDIA

Bigdata Analytics LAB Manual
B20EA0604
6th Semester

**SCHOOL OF COMPUTING AND INFORMATION
TECHNOLOGY**

Name	
SRN	
Branch	
Semester	
Section	
Academic Year	

Learning Objectives of the Course:

The objectives of this course are to:

1. Provide the knowledge of Map Reduce framework in solving problems related to big data.
2. Provide hands on experience on Hadoop environments.

Learning Outcomes of the Course:

Upon successful completion of the course, students should be able to:

1. Execute MapReduce programs on Hadoop and analyze the results.
2. Conduct some experiments on MapReduce.

CONTENTS		
Sl. No.	Lab Sessions	
	Installation and Configuration of Hadoop.	4
1	Write a Map Reduce program to solve the problem of word count for different file size.	5-8
2	Write a Map Reduce program to solve the problem of character count.	9-12
3	Write Map Reduce program to sort data by student names based on values.	13-14
4	Write Map Reduce program to process weather dataset and calculate average temperature.	15-17

INTRODUCTION:

Apache Hadoop is an open source framework that is used to efficiently store and process large datasets ranging in size from gigabytes to petabytes of data. Instead of using one large computer to store and process the data, Hadoop allows clustering multiple computers to analyze massive datasets in parallel more quickly.

MapReduce: A MapReduce job usually splits the input data-set into independent chunks which are processed by the map tasks in a completely parallel manner. The framework sorts the outputs of the maps, which are then input to the reduce tasks. Typically both the input and the output of the job are stored in a file-system.

Install Hadoop**Prerequisites:**

VMWare Workstation: it is used for installing the operating system on it.

OPERATING SYSTEM: You can install Hadoop on Linux-based operating systems. Ubuntu and CentOS are very commonly used. In this tutorial, we are using CentOS.

JAVA: You need to install the Java 8 package on your system.

HADOOP: You require Hadoop package.

1. Write a Map Reduce program to solve the problem of word count for Different file size.

WC Mapper.java

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
public class WC_Mapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text
value,OutputCollector<Text,IntWritable> output,
    Reporter reporter) throws IOException{
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()){
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

WC Reducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WC_Reducer extends MapReduceBase implements
Reducer<Text,IntWritable,Text,IntWritable> {
    public void reduce(Text key, Iterator<IntWritable>
values,OutputCollector<Text,IntWritable> output,
Reporter reporter) throws IOException {
        int sum=0;
        while (values.hasNext()) {
            sum+=values.next().get();
        }
        output.collect(key,new IntWritable(sum));
    }
}
```

WC Runner.java

```
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
```

```
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
public class WC_Runner {
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(WC_Runner.class);
        conf.setJobName("WordCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
        conf.setReducerClass(WC_Reducer.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf,new Path("input.txt"));
        FileOutputFormat.setOutputPath(conf,new Path("output"));
        JobClient.runJob(conf);
    }
}
```

INPUT:

Reva University is in Bangalore

Bangalore is in Karnataka

OUTPUT:

Bangalore,2

in, 2

is,2

Karnataka, 1

Reva, 1

University, 1

2. Write a Map Reduce program to solve the problem of Character count.**WC Mapper.java**

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WC_Mapper extends MapReduceBase implements
Mapper<LongWritable,Text,Text,IntWritable> {
    public void map(LongWritable key, Text
value,OutputCollector<Text,IntWritable> output,
        Reporter reporter) throws IOException{
        String line = value.toString();
        String tokenizer[] = line.split(" ");
        for(String SingleChar : tokenizer)
        {
            Text charKey = new Text(SingleChar);
            IntWritable One = new IntWritable(1);
            output.collect(charKey, One);
        }
    }
}
```

WC Reducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
public class WC_Reducer extends MapReduceBase implements
Reducer<Text,IntWritable,Text,IntWritable> {
    public void reduce(Text key, Iterator<IntWritable>
values,OutputCollector<Text,IntWritable> output,
Reporter reporter) throws IOException {
        int sum=0;
        while (values.hasNext()) {
            sum+=values.next().get();
        }
        output.collect(key,new IntWritable(sum));
    }
}
```

WC Runner.java

```
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
```

```
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
public class WC_Runner {
    public static void main(String[] args) throws IOException{
        JobConf conf = new JobConf(WC_Runner.class);
        conf.setJobName("CharCount");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(WC_Mapper.class);
        conf.setCombinerClass(WC_Reducer.class);
        conf.setReducerClass(WC_Reducer.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf,new Path("input.txt"));
        FileOutputFormat.setOutputPath(conf,new Path("output"));
        JobClient.runJob(conf);
    }
}
```

INPUT:

REVA UNIVERSITY

OUTPUT:

,1 //Number of lines

,1 //Number of white spaces

A,1

E,2

I,2

N,1

R,2

S,1

T,1

U,1

V,2

Y,1

3. Write a Map reduce program to sort data by student name based on values.

SortStudNames.java

Java project name: SortStudNames

Class name: SortStudNames.java

```
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class SortStudNames {
    public static class SortMapper extends Mapper <LongWritable, Text, Text,
    Text >
    {
        protected void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException {
            String[] token = value.toString().split(",");
            context.write(new Text(token[1]), new Text(token[0]+ "-" +token[1]));
        }
    }
    public static class SortReducer extends Reducer <Text, Text, NullWritable,
    Text>
    {
```

public void reduce(Text key, Iterable <Text> values, Context context) **throws**

IOException, InterruptedException {

for (Text details : values)

{

context.write(NullWritable.get(), details);

}

}

}

public static void main (String[] args) **throws** IOException,

InterruptedException, ClassNotFoundException

{

Configuration conf = **new** Configuration();

Job job = **new** Job(conf);

job.setJarByClass(SortStudNames.**class**);

job.setMapperClass(SortMapper.**class**);

job.setReducerClass(SortReducer.**class**);

job.setOutputKeyClass(Text.**class**);

job.setOutputValueClass(Text.**class**);

FileInputFormat.setInputPaths(job,**new** Path("input.csv"));

FileOutputFormat.setOutputPath(job,**new** Path("output"));

System.exit(job.waitForCompletion(**true**)? 0:1);

}

}

INPUT:

1001,Sahana

1003,Banu

1005,Deepa

1007,Akshay

OUTPUT:

1007,Akshay

1003,Banu

1005,Deepa

1001,Sahana

4. MapReduce program in Java for processing a weather dataset and calculate the average temperature of a particular year.**Project Name:WeatherMapReduce****Class Name:WeatherMapReduce**

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WeatherMapReduce {
    public static class WeatherMapper
    extends Mapper<LongWritable, Text, Text, DoubleWritable>{
        private final static DoubleWritable temperature = new DoubleWritable();
        private Text date = new Text();
        public void map(LongWritable key, Text value, Context context
        ) throws IOException, InterruptedException {
            String[] line = value.toString().split(",");
            if (line.length == 3) {
                date.set(line[0]);
                temperature.set(Double.parseDouble(line[2]));
                context.write(date, temperature);
            }
        }
    }
}
```

```
public static class WeatherReducer
extends Reducer<Text,DoubleWritable,Text,DoubleWritable> {
private DoubleWritable result = new DoubleWritable();
public void reduce(Text key, Iterable<DoubleWritable> values,
Context context
) throws IOException, InterruptedException {
double sum = 0;
int count = 0;
for (DoubleWritable val : values) {
sum += val.get();
count++;
}
double avg = sum / count;
result.set(avg);
context.write(key, result);
}
}

public static void main(String[] args) throws Exception {
Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "weather analysis");
job.setJarByClass(WeatherMapReduce.class);
job.setMapperClass(WeatherMapper.class);
job.setReducerClass(WeatherReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(DoubleWritable.class);
FileInputFormat.setInputPaths(job, new Path("input.txt"));
FileOutputFormat.setOutputPath(job, new Path("output"));
System.exit(job.waitForCompletion(true)?0 : 1);
}
}
```


INPUT:

Input.txt

2022-01-01,New York,32

2022-01-01,Los Angeles,55

2022-01-02,New York,30

2022-01-02,Los Angeles,58

2022-01-03,New York,35

2022-01-03,Los Angeles,60

2022-01-04,New York,28

2022-01-04,Los Angeles,56

2022-01-05,New York,33

2022-01-05,Los Angeles,61

2022-01-06,New York,29

2022-01-06,Los Angeles,59

2022-01-07,New York,36

2022-01-07,Los Angeles,57

Output:

2022-01-01 43.5

2022-01-02 44.0

2022-01-03 47.5

2022-01-04 42.0

2022-01-05 47.0

2022-01-06 44.0

2022-01-07 46.5